

Large Language Model Competitive Analysis

OpenAI | Anthropic | Google | Meta | Mistral

HyperAgent Research Division | April 2026

Executive Summary

The large language model (LLM) market has entered a new era of maturity and fierce competition as of Q2 2026. Five major providers dominate the landscape: OpenAI, Anthropic, Google, Meta, and Mistral. Each has carved distinct competitive positions through differentiated model families, pricing strategies, and platform capabilities.

Key Findings

1. OpenAI leads with GPT-5.4 family across three tiers (standard, mini, nano)
2. Google offers the broadest model portfolio with Gemini 3.x and aggressive free tiers
3. Anthropic differentiates through safety leadership and extended thinking with Claude 4
4. Meta continues to disrupt with fully open-source Llama 4 models
5. Mistral offers strong open-weight alternatives with European data sovereignty

This report provides a comprehensive comparison of each provider's latest model offerings, pricing structures, context window capabilities, and API features. Data was gathered directly from provider documentation and API pricing pages in April 2026.

Market Overview

The LLM API market has consolidated around a tiered pricing model: flagship reasoning models at premium prices, efficient mid-tier models for production workloads, and ultra-low-cost models for high-volume simple tasks. All five providers now offer some form of prompt caching, batch processing, and multimodal capabilities. The competitive battleground has shifted from raw intelligence to pricing efficiency, developer experience, and specialized capabilities such as agentic workflows, computer use, and code generation.

Context windows have expanded dramatically, with Google leading at 1M+ tokens for several models, and most flagship models supporting at least 200K tokens. Pricing per million tokens has dropped 60-80% compared to early 2024 levels, driven by architectural innovations and fierce competition.

3. Google (DeepMind)

Google offers the most extensive model portfolio through its Gemini API, spanning from the premium Gemini 3.1 Pro to the ultra-efficient 2.5 Flash-Lite. Google differentiates with generous free tiers, native multimodal capabilities, and deep integration with Google Search and Maps for grounding.

- Flagship: Gemini 3.1 Pro -- Best-in-class multimodal understanding and agentic capabilities
- Mid-tier: Gemini 3 Flash / 3.1 Flash-Lite -- Frontier performance at fraction of cost
- Economy: Gemini 2.5 Flash-Lite -- \$0.10/1M input, \$0.40/1M output
- Reasoning: Gemini 2.5 Flash -- Hybrid reasoning with 1M token context
- Specialized: Live API models, TTS, native image generation (Nano Banana)
- Processing tiers: Standard, Batch (-50%), Flex, Priority (+80%)
- Free tier available for most models with generous limits

Standard processing rates. All values in USD per 1 million tokens.

Provider	Model	Input	Cached	Output
OpenAI	GPT-5.4	\$2.50	\$0.25	\$15.00
Google	Gemini 3.1 Pro	\$2.00	\$0.20	\$12.00
Anthropic	Claude 4 Opus	\$15.00	\$3.75	\$75.00
Anthropic	Claude 4 Sonnet	\$3.00	\$0.75	\$15.00
Mistral	Mistral Large 3	\$2.00	--	\$6.00
Meta*	Llama 4 Maverick	\$0.20-0.50	--	\$0.40-1.20

* Meta models are open-source; prices shown are typical third-party API hosting rates.

Provider	Model	Input	Cached	Output
OpenAI	GPT-5.4 mini	\$0.75	\$0.075	\$4.50
Google	Gemini 3 Flash	\$0.50	\$0.05	\$3.00
Google	Gemini 3.1 Flash-Lite	\$0.25	\$0.025	\$1.50
Anthropic	Claude 4 Haiku	\$0.80	\$0.20	\$4.00
Mistral	Mistral Small 4	\$0.20	--	\$0.60
Meta*	Llama 4 Scout	\$0.10-0.25	--	\$0.20-0.60

Provider	Model	Input	Output
OpenAI	GPT-5.4 nano	\$0.20	\$1.25
Google	Gemini 2.5 Flash-Lite	\$0.10	\$0.40
Google	Gemini 2.5 Flash	\$0.30	\$2.50
Mistral	Ministral 3 (8B)	\$0.10	\$0.10
Mistral	Ministral 3 (3B)	\$0.04	\$0.10

Google Gemini 2.5 Flash-Lite at \$0.10/\$0.40 per 1M tokens offers the best price-performance among proprietary models for lightweight tasks.

Provider	Model	Context Window	Max Output
OpenAI	GPT-5.4	270K tokens	32K tokens
OpenAI	GPT-5.4 mini	270K tokens	32K tokens
OpenAI	GPT-5.4 nano	270K tokens	16K tokens
Google	Gemini 3.1 Pro	1M+ tokens	64K tokens
Google	Gemini 2.5 Pro	1M tokens	64K tokens
Google	Gemini 2.5 Flash	1M tokens	64K tokens
Google	Gemini 2.5 Flash-Lite	1M tokens	64K tokens
Anthropic	Claude 4 Opus	200K tokens	32K tokens
Anthropic	Claude 4 Sonnet	200K tokens	32K tokens
Anthropic	Claude 4 Haiku	200K tokens	16K tokens
Meta	Llama 4 Scout	10M tokens	Varies
Meta	Llama 4 Maverick	1M tokens	Varies
Mistral	Mistral Large 3	128K tokens	32K tokens
Mistral	Mistral Medium 3.1	128K tokens	32K tokens
Mistral	Mistral Small 4	128K tokens	32K tokens

Context Window Leaders

Meta Llama 4 Scout supports an extraordinary 10M token context window (practical throughput depends on hosting infrastructure). Google Gemini family consistently offers 1M token windows across most models. OpenAI and Anthropic prioritize quality over raw context size, with 200K-270K windows that include advanced prompt caching.

Prompt Caching Comparison

Provider	Cache Discount	Min Tokens	Cache TTL
OpenAI	90% off input	~1K tokens	5-10 min
Google	75-90% off input	Varies	Hourly billing
Anthropic	75% off input	1,024 tokens	5 min (auto-extend)
Mistral	Not yet available	--	--
Meta*	Provider-dependent	Varies	Varies

Feature	OpenAI	Google	Anthropic	Meta*	Mistral
Streaming	Yes	Yes	Yes	Yes	Yes
Function Calling	Yes	Yes	Yes	Yes	Yes
Structured JSON	Yes	Yes	Yes	Yes	Yes
Vision Input	Yes	Yes	Yes	Yes	Yes
Image Gen.	Yes	Yes	No	No	No
Audio Input	Yes	Yes	No	No	Yes
Audio/TTS	Yes	Yes	No	No	Yes
Video Input	Limited	Yes	No	Yes	No
Ext. Thinking	Yes	Yes	Yes	No	Yes
Computer Use	Yes	Yes	Yes	No	No
Web Search	Yes	Yes	Yes	No	Yes
Code Exec.	Yes	Yes	Yes	No	No
Batch API	-50%	-50%	-50%	Varies	Yes
Caching	90% off	75-90%	75% off	Varies	No
Fine-tuning	SFT/DPO/RFT	Yes	Yes	Open wt.	Yes
Embeddings	Yes	Yes	Voyage	No	Yes
Moderation	Yes	Yes	Yes	No	Yes
OAI-compat.	--	Yes	Yes	Yes	Yes

* Meta provides open-source models; API features depend on the hosting provider.

Processing Tier	OpenAI	Google	Anthropic
Standard	Base price	Base price	Base price
Batch (async, 24h)	-50%	-50%	-50%
Flex (lower priority)	Available	Same as Batch	--
Priority (SLA)	Scale Tier	+80%	--
Free Tier	No	Yes (generous)	Limited
Data Residency	+10%	Via Vertex AI	--

Consideration	OpenAI	Google	Anthropic	Meta	Mistral
SOC 2	Yes	Yes	Yes	N/A	Yes
Data Residency	+10%	Vertex AI	Limited	Self-host	EU Native
SLA Guarantees	Scale Tier	Enterprise	Available	Self-host	Enterprise
Volume Discount	Reserved	Yes	Yes	N/A	Yes
On-Premise	No	GKE	No	Yes	Yes

GDPR Focus	Standard	Standard	Standard	Self-host	EU-first

Use Case	Recommended	Runner-Up	Rationale
Complex reasoning	OpenAI GPT-5.4	Claude 4 Opus	Best cost/capability balance
High-volume prod.	Gemini 2.5 Flash-Lite	GPT-5.4 nano	Lowest cost per token
Code generation	GPT-5.4 mini	Devstral 2	Strong coding + agents
Document analysis	Gemini 3.1 Pro	Claude 4 Sonnet	1M context for long docs
Cost-sensitive	Llama 4 (self-host)	Google free tier	Zero per-token cost
EU compliance	Mistral Large 3	Llama 4 self-host	Native EU residency
Real-time voice	GPT-realtime-1.5	Gemini Live	Most mature voice API
Safety-critical	Claude 4 Opus	GPT-5.4	Constitutional AI

Full multimodal	Gemini 3.1 Pro	GPT-5.4	Native text+img+vid+audio
Market Outlook			
The LLM market is evolving rapidly toward agentic AI, where models autonomously use tools, browse the web, write and execute code, and orchestrate multi-step workflows. All five providers are investing heavily in this direction. We expect pricing pressure to continue, with 30-50% further reductions likely by end of 2026 as architectural efficiency improves and competition intensifies. Key trends to watch: (1) the convergence of reasoning and traditional models into hybrid architectures, (2) native multimodal generation becoming table stakes, (3) the growing importance of tool ecosystems over raw model intelligence, and (4) open-source models closing the gap with proprietary offerings.			

Methodology Note

Pricing data for OpenAI and Google was sourced directly from their official API pricing pages. Anthropic and Meta pricing reflects publicly available documentation and third-party benchmarks. All prices are USD, standard processing rates unless noted. Verify current rates at each provider's official pricing page before making procurement decisions.

Report prepared by HyperAgent Research Division -- April 2026
For questions or updates, visit github.com/hyperlight-dev/hyperagent